

# Advanced Mathematical, ML, and Infrastructure Refinements for NEPSE Quant System

**Purpose:** This addendum captures critical last-mile refinements that materially improve robustness, calibration, and survivability of a probabilistic trading system in a thin, regime-driven emerging market. This document is complementary to the main blueprint and assumes full familiarity with it.

---

## 1. Mathematical Refinements for Thin-Market Dynamics

### 1.1 Asymmetric Power-Law Market Impact

In illiquid markets, price impact is asymmetric: sell pressure in stress regimes produces larger and faster price dislocations than buy pressure. Model impact separately for buys and sells.

Let trade size be  $q_t$  (signed, positive = buy, negative = sell), ADV be average daily volume.

Define asymmetric impact:

$$\text{Impact}(q_t) = \begin{cases} k_1^+ \left( \frac{q_t}{\text{ADV}} \right)^{\alpha^+} + k_2^+ \left( \frac{q_t}{\text{ADV}} \right)^2 & q_t > 0 \\ k_1^- \left( \frac{|q_t|}{\text{ADV}} \right)^{\alpha^-} + k_2^- \left( \frac{|q_t|}{\text{ADV}} \right)^2 & q_t < 0 \end{cases}$$

Where typically:  $-\alpha^- > \alpha^+$   $-k^- > k^+$  during stress regimes

**Regime-conditioned parameters:**

$$(k_1^\pm, \alpha^\pm) = f(s_t)$$

Calibrate parameters using regime-filtered execution data via rolling nonlinear least squares or recursive least squares (RLS).

---

### 1.2 Hierarchical Bayesian Shrinkage for Thin Stocks

For stocks with sparse trading, estimate parameters via partial pooling.

Example: volatility or impact coefficient  $\theta_i$  for stock  $i$ :

$$\theta_i \sim \mathcal{N}(\mu_{\text{NEPSE}}, \tau^2)$$

Observed estimate:

$$\hat{\theta}_i \mid \theta_i \sim \mathcal{N}(\theta_i, \sigma_i^2)$$

Posterior mean:

$$\mathbb{E}[\theta_i \mid \hat{\theta}_i] = w_i \hat{\theta}_i + (1 - w_i) \mu_{\text{NEPSE}}, \quad w_i = \frac{\tau^2}{\tau^2 + \sigma_i^2}$$

As data accumulates ( $\sigma_i^2 \downarrow$ ), estimates naturally de-shrink.

Use for: - Impact parameters - Volatility persistence - Regime transition probabilities

---

### 1.3 CVaR-Constrained Portfolio Optimization

Replace stop-loss logic with tail-aware optimization.

Let portfolio returns be  $R_p = w^\top R$ .

Define CVaR at level  $\alpha$ :

$$\text{CVaR}_\alpha(R_p) = \mathbb{E}[R_p \mid R_p \leq \text{VaR}_\alpha]$$

Optimization problem:

$$\max_w \mathbb{E}[R_p] - \lambda \cdot \text{CVaR}_\alpha(R_p)$$

Subject to: - Sector exposure caps - Liquidity constraints - Gross and net exposure limits

Solved via linear programming using scenario simulation from regime-weighted predictive distributions.

---

## 2. Programming & Infrastructure Enhancements

### 2.1 Stateful LOB Matching Engine (Numba/Cython)

The LOB simulator must preserve event ordering and queue state.

Key state variables: - Queue depth at each price level - Order arrival timestamps - Agent queue position

Fill probability model:

$$\mathbb{P}(\text{fill within } \Delta t) = 1 - \exp\left(-\frac{V_{\text{ahead}}}{\lambda_{\text{market}}}\right)$$

Where: -  $V_{\text{ahead}}$  = volume ahead in queue -  $\lambda_{\text{market}}$  = arrival rate of opposing market orders

Numba/Cython used to process  $10^6+$  events/sec in backtests and simulation.

---

### 2.2 Feature Parity Auditing (Offline vs Live)

Every feature  $x_t$  must satisfy:

$$x_t^{\text{train}} \equiv x_t^{\text{live}}$$

Implementation: - Deterministic rolling windows - Fixed seed ordering - Numerical hash (xxhash / SHA256) on feature vectors

$$H(x_t^{\text{train}}) = H(x_t^{\text{live}})$$

Mismatch triggers hard fail and trading halt.

---

## 2.3 Hardware Timestamping & Clock Sync

Queue-sensitive strategies require sub-millisecond sync.

- Use **PTP (IEEE 1588)** with hardware NIC timestamping
- Target clock skew: < 50 microseconds

Execution engine timestamps: - Market data arrival - Feature snapshot - Order submission - Exchange ACK

Latency budget must be explicitly tracked and logged.

---

## 3. Machine Learning Enhancements

### 3.1 Regime-Conditional Feature Importance

Estimate importance conditional on regime:

$$I_j^{(k)} = \mathbb{E}[\Delta L \mid s_t = k]$$

Where  $\Delta L$  is loss increase under permutation of feature  $j$ .

Findings often show: - Ownership / promoter features dominate in mean-reversion regimes - Liquidity and volatility dominate in stress regimes

---

### 3.2 Proper Probabilistic Scoring (CRPS)

Train ML models to optimize distributional accuracy.

Continuous Ranked Probability Score:

$$\text{CRPS}(F, y) = \int_{-\infty}^{\infty} (F(z) - 1\{z \geq y\})^2 dz$$

CRPS penalizes: - Overconfident wrong predictions - Underconfident correct predictions

Used as primary objective for: - Quantile regression - NGBoost / probabilistic boosting

---

### 3.3 Ensemble Anomaly Detection

Combine: 1. Deterministic rules (price jumps, zero volume) 2. Isolation Forest trained on clean historical states

Final anomaly score:

$$A_t = \max(A_t^{\text{rules}}, A_t^{\text{IF}})$$

If  $A_t > \tau$ , data is quarantined and no signal is generated.

---

## 4. Monitoring & Drift Detection

### 4.1 Feature Population Stability Index (PSI)

For feature  $x$ :

$$\text{PSI} = \sum_i (p_i - q_i) \log \left( \frac{p_i}{q_i} \right)$$

- $p_i$ : training distribution
- $q_i$ : live distribution

PSI > 0.2 triggers investigation; > 0.3 triggers model retraining or shutdown.

---

## 5. Summary of Enhancements

Category	Enhancement	Benefit
Math	Asymmetric Impact	Captures sell-side fragility
Math	Hierarchical Shrinkage	Stabilizes thin-stock estimates
Risk	CVaR Optimization	Controls tail losses
Infra	Numba/Cython LOB	Realistic execution simulation
Infra	PTP Sync	Queue-position accuracy
ML	CRPS Objective	Better probabilistic calibration
ML	Regime Feature Importance	Structural interpretability
Ops	PSI Monitoring	Early regime drift detection

---

*End of addendum.*