

# Production-grade Probabilistic Quantitative + ML Trading System for an Illiquid Emerging Market (NEPSE)

**Audience:** Quantitative researcher with strong math and Python skills. Institutional depth; assume no prior conversation context.

---

## Table of contents

1. Problem Definition
  2. Data Preparation & Stationarity
  3. Market Regime Modeling
  4. Volatility Modeling
  5. Liquidity & Market Impact
  6. Decomposition: Linear Structure vs Nonlinear Residuals
  7. Machine Learning Layer
  8. Forecast Combination
  9. Signal Extraction (MOST IMPORTANT)
  10. Position Sizing & Risk Control
  11. Backtesting & Evaluation
  12. Profitability Reality Check
  13. Scalability & Extensions
- 

## 1. Problem Definition

**What is being predicted.**

Define the *random variable* explicitly. Let  $P_t$  be the mid-price at time  $t$ . For a fixed horizon  $h$  (e.g., 1 day, 5 days), define the log-return:

$$R_{t,h} \equiv \log \frac{P_{t+h}}{P_t}.$$

The system must produce a **probabilistic forecast** of  $R_{t,h}$  conditional on the filtration  $\mathcal{F}_t$  (all information available at time  $t$ ).

$$\text{Goal: } \hat{F}_{t,h}(x) \approx \mathbb{P}(R_{t,h} \leq x \mid \mathcal{F}_t),$$

or equivalently provide the conditional density  $f_{t,h}(x)$  and moments (mean, variance, higher moments or quantiles).

### **Why price prediction is wrong; why returns are used.**

- Price is a nonstationary process with multiplicative scale and trending level shifts. Predicting price level is dominated by drift and scale effects; such predictions are poor for execution. Returns (log-returns) remove multiplicative scale and allow modeling of stationary increments (or detect when stationarity fails).
- A point prediction  $\mathbb{E}[R_{t,h}]$  is insufficient; execution decisions require the *distribution* of outcomes (tail risk, skewness, kurtosis) and probability mass around the bid-ask and slippage thresholds.

### **Mathematical definition of the target.**

Full target: conditional distribution  $F_{t,h}$  or a parametric family  $\mathcal{D}(\theta_{t,h})$  where parameters  $\theta_{t,h}$  are predicted from features. Example:  $R_{t,h} | \mathcal{F}_t \sim t(\mu_{t,h}, \sigma_{t,h}, \nu_{t,h})$  where  $\mu, \sigma, \nu$  are time-varying.

---

## **2. Data Preparation & Stationarity**

### **Log returns.**

Definition (single-step):

$$r_t = \log P_t - \log P_{t-1} = \log \left( \frac{P_t}{P_{t-1}} \right).$$

Multi-step horizon  $h$ :

$$R_{t,h} = \log \left( \frac{P_{t+h}}{P_t} \right) = \sum_{i=1}^h r_{t+i}.$$

### **Why required.**

- Additivity across horizons.
- Stabilizes multiplicative effects; aligns with continuous-time returns  $d \log P$ .
- Enables use of standard time-series tools that assume stationarity (or piecewise stationarity).

### **Stationarity assumptions and consequences of violating them.**

Assume:  $r_t$  is weakly stationary conditional on regimes: constant mean and autocovariance structure within a regime. In emerging markets, stationarity often fails globally due to regime shifts (policy, market microstructure events) and illiquidity spikes; consequences:

- Model coefficients become time-varying—necessitating regime modeling or robust online updating.
- Estimators become biased; standard errors incorrect.
- Backtest Sharpe and VaR estimates are unreliable.

### **Python: compute returns and quick stationarity checks**

```

import numpy as np
import pandas as pd
from statsmodels.tsa.stattools import adfuller

# price series: pandas Series indexed by timestamp
def log_returns(price: pd.Series) -> pd.Series:
    return np.log(price).diff().dropna()

# ADF test for stationarity
def adf_test(series: pd.Series):
    stat, pvalue, _, _, crit_vals, _ = adfuller(series.dropna())
    return dict(stat=stat, pvalue=pvalue, crit=crit_vals)

# Example
# prices = pd.read_csv('prices.csv', index_col=0, parse_dates=True)['mid']
# r = log_returns(prices)
# print(adf_test(r))

```

Notes: do not blindly accept ADF in illiquid data—thin trading produces zero returns and spurious unit roots.

---

### 3. Market Regime Modeling

#### Model: Markov Regime-Switching Autoregressive (MS-AR)

Full (econometric) equation for  $r_t$  with  $K$  regimes and AR(p) dynamics within each regime:

Let  $s_t \in \{1, \dots, K\}$  be an unobserved Markov chain with transition matrix  $Q$  where  $q_{ij} = \mathbb{P}(s_t = j | s_{t-1} = i)$ . Conditional on regime  $s_t = k$ :

$$r_t = \mu^{(k)} + \sum_{j=1}^p \phi_j^{(k)} r_{t-j} + \sigma^{(k)} \varepsilon_t^{(k)}, \quad \varepsilon_t^{(k)} \stackrel{iid}{\sim} D(0, 1),$$

where  $D$  is a standardized distribution (Gaussian or t). Transition dynamics:

$$\mathbb{P}(s_t = j | s_{t-1} = i) = q_{ij}.$$

#### Explanation of regimes

- *Trend* regime:  $\phi_1^{(k)} > 0$ , positive autocorrelation; persistent directional moves.
- *Mean-reversion* regime: negative autocorrelation,  $\sum \phi_j^{(k)}$  pulling back to mean.
- *Stress* regime: high  $\sigma^{(k)}$ , fat tails, potentially skewed disturbances.

#### Why regime-awareness is mandatory in emerging markets

- Nonstationarity is dominated by infrequent but structural regime switches. Liquidity, policy moves, and concentrated ownership create abrupt transitions.

- Trading signals that ignore regimes will be miscalibrated for volatility and autocorrelation, causing large P&L drawdowns when regime flips.

**Python implementation using statsmodels**

```

import numpy as np
import pandas as pd
from statsmodels.tsa.regime_switching.markov_autoregression import
MarkovAutoregression

# r: pandas Series of returns
# Fit MS-AR with K regimes and order p
def fit_ms_ar(r: pd.Series, k_regimes=3, order=1, trend='c'):
    model = MarkovAutoregression(r, k_regimes=k_regimes, order=order,
trend=trend)
    res = model.fit(em_iter=50, search_reps=5, method='nm')
    # res.summary() gives coefficients and smoothed probabilities
    smooth_probs = res.smoothed_marginal_probabilities
    return res, smooth_probs

# Example
# res, probs = fit_ms_ar(r, k_regimes=3, order=1)
# probs.tail()

```

Interpretation: use `res.predict()` and `res.smoothed_marginal_probabilities` to obtain regime probabilities  $\mathbb{P}(s_t = k | \mathcal{F}_t)$ . These probabilities feed downstream features and dynamic weighting.

## 4. Volatility Modeling

**Choice:** EGARCH (asymmetric, models log-variance) or GJR-GARCH (captures leverage effect). Use heavy-tailed innovations (Student-t) to model fat tails.

**EGARCH(1,1) formulation.** Let  $\sigma_t^2$  be conditional variance. Using EGARCH(p,q) with p=q=1:

$$\log \sigma_t^2 = \omega + \beta \log \sigma_{t-1}^2 + \alpha \left( \frac{|\varepsilon_{t-1}|}{\sigma_{t-1}} - E \left[ \frac{|\varepsilon_{t-1}|}{\sigma_{t-1}} \right] \right) + \gamma \frac{\varepsilon_{t-1}}{\sigma_{t-1}},$$

where  $\varepsilon_t = r_t - \mu_t$  and  $\gamma$  captures asymmetry (leverage).

**Why volatility is a forecast, not noise.**

- Volatility clusters and is predictable conditional on past innovation magnitudes, liquidity, and regime probabilities; treating it as noise ignores predictable risk and mis-scales position size.

**Fat tails and leverage.**

- Use Student-t or generalized error distribution for  $\varepsilon_t$  to capture tails.

- Include asymmetric terms (EGARCH / GJR) to let negative shocks impact future variance more than positive shocks.

**Python implementation using arch**

```
from arch import arch_model

# r: pandas Series of returns
# Fit EGARCH(1,1) with t-distribution
def fit_egarch(r: pd.Series):
    am = arch_model(r*100, mean='Constant', vol='EGARCH', p=1, q=1,
dist='StudentsT')
    res = am.fit(disp='off')
    return res

# Fit GJR-GARCH(1,1)
def fit_gjr(r: pd.Series):
    am = arch_model(r*100, mean='Constant', vol='GARCH', p=1, o=1, q=1,
dist='StudentsT')
    res = am.fit(disp='off')
    return res

# Example
# eg_res = fit_egarch(r)
# print(eg_res.summary())
```

Scale returns (multiply by 100) for numerical stability in arch package. Use the forecasted  $\hat{\sigma}_{t+h}$  to standardize residuals and compute probabilistic forecasts.

---

## 5. Liquidity & Market Impact

**Amihud illiquidity measure.** For asset  $i$ , over a window of  $d$  days:

$$\text{ILLIQ}_{i,t} = \frac{1}{d} \sum_{\tau=t-d+1}^t \frac{|R_\tau|}{\text{Volume}_\tau},$$

units: return-per-unit-volume. Interpretation: higher values = more price move for given volume (illiquid).

**Why ignoring liquidity makes backtests fake.**

- Assumed execution at mid-price and unlimited size is unrealistic. In illiquid markets, market impact and crossing spreads dominate small predicted edges. Backtests that ignore liquidity overstate returns and understate drawdowns.

**Liquidity as a penalty term.**

In expected return calculation, subtract impact/cost:

$$\text{AdjustedEdge}_t = \mathbb{E}[R_{t,h} | \mathcal{F}_t] - c_0 - c_1 \cdot \text{ILLIQ}_t - c_2 \cdot \text{Spread}_t,$$

where  $c_1, c_2$  are calibration parameters (estimated from microstructure or backtests). Alternatively, incorporate liquidity into position cap:

$$\text{TradeSize}_{\max} = \min \left( \text{VaR-based cap}, \kappa \cdot \frac{\text{ADV}}{\text{market impact factor}} \right).$$


---

## 6. Decomposition: Linear Structure vs Nonlinear Residuals

**Model decomposition.** Decompose returns as:

$$r_t = L_t + \varepsilon_t, \quad L_t = \mathbf{x}_t^\top \boldsymbol{\beta}_t$$

where  $L_t$  is the structured (linear) component, e.g., AR terms, regime-dependent means, and observable exogenous linear predictors; and  $\varepsilon_t$  is the residual. The residual is what the ML layer should model.

**Mathematical definition of residuals.**

Given estimated linear predictor  $\hat{L}_t$ , residual:

$$\hat{\varepsilon}_t = r_t - \hat{L}_t.$$

Standardize residuals by forecasted volatility  $\hat{\sigma}_t$ :

$$\tilde{\varepsilon}_t = \frac{\hat{\varepsilon}_t}{\hat{\sigma}_t}.$$

**Why ML should only model residuals.**

- Linear structure captures predictable, interpretable dynamics with low variance of estimation. If ML models raw returns, it risks re-learning linear relations and overfitting noise.
  - ML focuses capacity on nonlinear interactions in standardized residuals where extra explanatory power might exist.
- 

## 7. Machine Learning Layer

**Model choice:** LightGBM or XGBoost for gradient-boosted trees as primary; neural nets only as an extension with strict regularization and temporal CV.

**Feature set (example):**

- Lagged returns:  $r_{t-1}, \dots, r_{t-p}$
- Lagged absolute returns:  $|r_{t-1}|, \dots$

- Volatility forecasts:  $\hat{\sigma}_t, \hat{\sigma}_{t+1}$
- Regime probabilities:  $\mathbb{P}(s_t = k)$  for each  $k$
- Liquidity measures: Amihud, ADV, spread
- Orderbook features (if available): top-of-book imbalance, depth ratios
- Calendar features: day-of-week, month
- Cross-sectional features (when moving to multiple stocks): market cap, float, concentration

### Training target.

Train on standardized residuals or directly on tail-probabilities. Preferred target:

$$Y_t = \mathbb{E} \left[ \frac{R_{t,h}}{\hat{\sigma}_{t,h}} \mid \mathcal{F}_t \right] \approx \mathbb{E}[\tilde{\varepsilon}_{t,h} \mid \mathcal{F}_t].$$

For probabilistic outputs, train models to predict multiple quantiles or distribution parameters (e.g., gradient boosting for quantile regression or use NGBoost for probabilistic boosting).

### Avoiding leakage.

- Use strictly causal features (no future data).
- Train/validate via expanding window forward chaining (no random shuffle).
- When features use rolling statistics, ensure only past window is used.
- Use purging to remove overlapping labels if horizons overlap.

### Python example (LightGBM) — training on residuals

```
import lightgbm as lgb
from sklearn.model_selection import TimeSeriesSplit

# X: DataFrame of features indexed by time. y: target = standardized
# residuals for horizon h

def train_lgb(X, y, n_splits=5):
    tss = TimeSeriesSplit(n_splits=n_splits)
    models = []
    for train_idx, val_idx in tss.split(X):
        X_tr, X_val = X.iloc[train_idx], X.iloc[val_idx]
        y_tr, y_val = y.iloc[train_idx], y.iloc[val_idx]
        dtrain = lgb.Dataset(X_tr, y_tr)
        dval = lgb.Dataset(X_val, y_val)
        params = {
            'objective': 'regression',
            'metric': 'l2',
            'learning_rate': 0.01,
            'num_leaves': 31,
            'min_data_in_leaf': 100,
            'lambda_l1': 1.0,
            'lambda_l2': 1.0,
            'verbose': -1
        }
        bst = lgb.train(params, dtrain, valid_sets=[dval],
```

```

    early_stopping_rounds=100, num_boost_round=2000)
        models.append(bst)
    return models

# Predict: ensemble average of models

```

Notes: use strong regularization (min\_data\_in\_leaf, lambda L1/L2) and low learning rate. Use feature importance diagnostics and SHAP for interpretability.

---

## 8. Forecast Combination

**Combine linear forecast and ML residual forecast.**

Let:\

- $L_{t,h}$ : linear forecast for horizon  $h$  (from MS-AR predicted mean) -  $\hat{\varepsilon}_{t,h}^{ML}$ : ML forecast for standardized residual -  $\hat{\sigma}_{t,h}$ : volatility forecast

Construct conditional mean forecast:

$$\hat{\mu}_{t,h} = L_{t,h} + \hat{\sigma}_{t,h} \cdot \hat{\varepsilon}_{t,h}^{ML}.$$

For full predictive distribution, combine parametric residual distribution (e.g., Student-t with fitted  $\hat{\nu}$ ) with estimated mean and scale:

$$R_{t,h} \mid \mathcal{F}_t \sim t(\mu = \hat{\mu}_{t,h}, \sigma = \hat{\sigma}_{t,h}, \nu = \hat{\nu}_{t,h}).$$

**Weighting logic.**

- Fixed weighting: choose scalar  $\alpha \in [0, 1]$  and set final mean =  $(1 - \alpha)L + \alpha(L + \sigma\varepsilon^{ML})$  — equivalently trust ML fraction  $\alpha$ .
- Dynamic weighting: let  $\alpha_t$  be inverse of expected conditional MSE estimated from CV; compute rolling out-of-sample RMSEs for linear and ML components and set weights proportional to inverse RMSE.

Formally, let  $MSE_t^{lin}, MSE_t^{ml}$  be forward-looking CV errors; set

$$\alpha_t = \frac{MSE_t^{lin}}{MSE_t^{lin} + MSE_t^{ml}}.$$

**Final predictive distribution.** Use mixture or single parametric with fat tails. If residuals non-Gaussian, use a parametric family for tails or a mixture-of-experts combining regime-specific distributions.

---

## 9. Signal Extraction (MOST IMPORTANT)

**Predictions → Signals (not trades).**

Define the **edge** as expected net return per unit of traded risk after costs and impact.

Mathematical edge for a proposed trade of sign  $s \in \{+1, -1\}$  and intended notional fraction  $x$  of portfolio:

$$\text{Edge}(x) = x \cdot \hat{\mu}_{t,h} - \text{Cost}(x) - \text{Impact}(x) - \lambda_{\text{liq}} \cdot f(\text{ILLIQ}_t, x),$$

where  $\hat{\mu}_{t,h}$  is forecasted return,  $\text{Cost}(x)$  includes explicit fees and half-spread,  $\text{Impact}(x)$  is expected adverse price movement from executing size  $x$ , and  $\lambda_{\text{liq}}$  penalizes illiquidity.

### Why accuracy % is meaningless.

- Classification accuracy ignores calibration and economic significance. A 60% accuracy with tiny edges may be useless; a 40% accuracy with large positive expectancy per trade (asymmetric payoff) could be profitable. Expected value per trade and risk per trade matter.

### Threshold logic.

Set a threshold  $\tau$  on Edge/TradeRisk where TradeRisk is e.g., forecast volatility scaled by size.

Formulate z-score of edge:

$$Z_t = \frac{\hat{\mu}_{t,h} - \text{costs}}{\hat{\sigma}_{t,h}}.$$

Trade only if  $|Z_t| > \tau$ . Choose  $\tau$  to control expected turnover and ensure positive expected P&L after costs. Example heuristics:

- Calibrate  $\tau$  via rolling OOS backtest to maximize realized IR under transaction cost budget.
- Use separate thresholds per regime and liquidity bucket.

### Mathematical threshold selection (example).

Maximize expected utility under linear utility with risk-aversion parameter  $\gamma$ : choose trade size  $x$  solving

$$\max_x x\hat{\mu} - x^2 \frac{\hat{\sigma}^2}{2\gamma} - \text{Cost}(x).$$

First-order condition yields threshold on  $\hat{\mu}/\hat{\sigma}$ . Solve for minimal  $\hat{\mu}$  that leads to nonzero  $x$ .

### Python: turning predictive distribution into signals

```
import numpy as np

# inputs per asset/time
mu_hat = ...          # forecasted mean return (horizon h)
sigma_hat = ...         # forecasted vol
illiq = ...             # Amihud
spread = ...            # half-spread
fee = 0.0005

# cost model approx: linear + quadratic impact
def cost_estimate(x, adv, illiq, spread, fee, k1=0.1, k2=0.3):
```

```

# x: fraction of capital to trade
# adv: average daily volume (fraction of float)
market_impact = k1 * (x/adv) + k2 * (x/adv)**2
explicit = spread + fee
liq_penalty = illiq * 1e3
return explicit + market_impact + liq_penalty

# Signal generation
def make_signal(mu_hat, sigma_hat, illiq, spread, adv, fee=0.0005, tau=0.5):
    # z-score
    z = mu_hat / sigma_hat
    # naive desired fraction from Kelly-like
    f_kelly = mu_hat / (sigma_hat**2 + 1e-9)
    # compute costs at candidate f
    cost = cost_estimate(abs(f_kelly), adv, illiq, spread, fee)
    adjusted_edge = mu_hat - cost
    z_adj = adjusted_edge / sigma_hat
    if abs(z_adj) > tau:
        sign = np.sign(z_adj)
        size = np.clip(0.5 * f_kelly, -1.0, 1.0) # conservative scale
        return dict(signal=sign, size=size, z=z_adj)
    return dict(signal=0, size=0, z=z_adj)

# Example usage in vectorized backtest loop

```

Notes: signals must be stored as time-stamped records, with regime, liquidity bucket, and expected cost attached for post-trade analysis.

## 10. Position Sizing & Risk Control

**Volatility targeting and Kelly-style sizing.**

- *Volatility targeting*: choose position size such that portfolio ex-ante volatility equals target  $\sigma^*$ :

$$w_t = \frac{\sigma^*}{\hat{\sigma}_t} \cdot \frac{\hat{\mu}_t}{|\hat{\mu}_t|},$$

but volatility targeting alone ignores expected return magnitude. Better: fractional Kelly:

$$f_t = \eta \cdot \frac{\hat{\mu}_t}{\hat{\sigma}_t^2}, \quad \eta \in (0, 1) \text{ for shrinkage}$$

cap and liquidity adjustment:

$$f_t^{adj} = f_t \cdot \frac{1}{1 + \zeta \cdot ILLIQ_t}, \quad |f_t^{adj}| \leq f_{\max}$$

### **Drawdown containment logic.**

- Stop-loss by portfolio drawdown: reduce leverage stepwise after drawdown thresholds (e.g., reduce by 25% after 5% drawdown, 50% after 10%).
- Volatility / tail risk triggers: if realized vol > k \* forecast vol or stress regime probability > 0.6, move to cash or reduce size.
- Hard capital allocation caps by single-asset exposure and daily turnover limits.

**Risk overlays.** Implement delta limits, stop-losses, sector-level caps, correlation stress tests, and daily max loss per trader/strategy.

---

## **11. Backtesting & Evaluation**

### **Metrics that matter.**

- Information Ratio / Sharpe (with robust standard error).
- Maximum drawdown and duration.
- Turnover and transaction costs.
- Liquidity consumption and market impact realized vs modeled.
- Tail risk metrics: 95% Var, expected shortfall.
- Profit per trade, win size vs loss size distributions.

### **Why win-rate is misleading.**

Win-rate ignores magnitude — large losses can wipe out many small wins. Expected return per trade and distributional skewness/kurtosis matter more.

### **Common backtest lies and mitigations.**

- *Lookahead / Label leakage*: avoid by strictly causal features and time-series CV. Use purging when labels overlap.
- *Survivorship bias*: use full universe including delisted stocks.
- *Transaction cost underestimation*: model spread, fees, and price impact calibrated to real fills.
- *Ignoring execution latency and queueing*: in illiquid markets, model order matching and partial fills.

### **Backtest skeleton (walk-forward)**

```
# Pseudocode: expanding-window walk-forward
window_train = 252*3 # 3 years
step = 21 # re-train monthly
for start in range(0, n - window_train - test_size, step):
    train_slice = slice(start, start+window_train)
    test_slice = slice(start+window_train, start+window_train+step)
    # fit linear / regime / vol / ml on train
    # produce forecasts for test
    # simulate execution with cost model and sizing
```

```
# record PnL, turnover, drawdown  
# aggregate OOS metrics
```

## 12. Profitability Reality Check

### What determines profitability.

- True, persistent information content in features that is not already arbitraged away.
- Low friction or predictable liquidity windows to execute without large impact.
- Robust risk management and regime-aware reduction of exposure during structural changes.

### Expected statistical characteristics (ranges, not promises).

For an alpha on an illiquid emerging index:

- Annualized expected return *alpha* may be small (1–10% excess depending on capital and costs).
- Expect high turnover and variable realized Sharpe; realistic target IR often 0.3–1.0 for mid-frequency strategies after costs in such markets.

### Why most such systems fail.

- Overfitting to noise, failure to model costs and liquidity, structural breaks, insufficient capacity, and organizational execution weaknesses.

## 13. Scalability & Extensions

### From index to stocks.

- Move from single-series modeling to panel models. Use cross-sectional ML that borrows strength across assets (multi-task learning) but account for differing liquidity and idiosyncratic microstructure.
- Add hierarchical risk models: factor exposures, sector constraints.

### Increasing ML complexity safely.

- Use holdout sets, nested CV, and model monitoring. Increase complexity only if OOS performance improves in a stable way across regimes.
- Regularize heavily and use calibration targets (probabilistic scoring rules like CRPS or negative log-likelihood) rather than raw RMSE.

### What breaks first as capital increases.

- Liquidity: realized market impact grows nonlinearly with size.
- Slippage and opportunity cost from inability to deploy full size.
- Strategy crowding leading to alpha decay.

# Appendix: Compact formulas and code snippets

## 1) Target & density

$$R_{t,h} = \sum_{i=1}^h r_{t+i}, \quad r_t = \log P_t - \log P_{t-1}.$$

Parametric forecast:  $R_{t,h} \mid \mathcal{F}_t \sim t(\mu_{t,h}, \sigma_{t,h}, \nu_{t,h})$ .

## 2) Amihud

$$\text{ILLIQ}_t = \frac{1}{d} \sum_{\tau=t-d+1}^t \frac{|R_\tau|}{\text{Volume}_\tau}.$$

## 3) Kelly fraction (fractional)

$$f_t = \eta \cdot \frac{\mu_t}{\sigma_t^2}, \quad 0 < \eta < 1.$$

## 4) Simple cost model

$$\text{Cost}(x) = \text{spread} + c_1 \frac{x}{\text{ADV}} + c_2 \left( \frac{x}{\text{ADV}} \right)^2.$$

---

# Final notes on productionization

- Implement modular architecture: data ingestion layer (resilient to missing ticks), feature store, model training pipelines (airflow or prefect), model registry, real-time scoring service, and execution engine with simulator and real broker connectors.
- Monitoring: model drift detection, PnL attribution, regime-probability monitoring, and automated rollback.
- Governance: documented assumptions, parameter provenance, and independent validation of cost/impact models.

---

*End of document.*